ChatGPT

UK compute capacity for AI inference – mapping of public-sector-accessible infrastructure (2026)

Introduction

The UK government plans to grow its AI research and innovation compute capacity twenty-fold to **420 AI-exaflops by 2030** through a balanced mix of purpose-built supercomputers and cloud-based resources [1]. The **AI Research Resource (AIRR)** is central to this plan, providing large-scale GPU and specialised inference systems that can be accessed by academics, start-ups and public-sector organisations. Beyond the AIRR, the UK hosts a rich ecosystem of Tier-0, Tier-1 and Tier-2 high-performance computing (HPC) centres, data-science clouds and research facilities that support AI inference. The map below illustrates the distribution of major compute centres across the UK.

The sections below provide a comprehensive inventory of publicly funded compute resources and explain how they can be leveraged for AI inference in the public sector. Tables summarise hardware and access information; prose sections discuss applications and examples.

1 AI Research Resource (AIRR) and national AI supercomputers

1.1 Isambard-AI (University of Bristol – National Composites Centre)

- **Purpose & status.** Isambard-AI is the UK's flagship AI supercomputer and forms half of the AIRR. It opened in July 2025 at the National Composites Centre in Bristol. The facility cost **£225 million** and is

run by the University of Bristol and Hewlett Packard Enterprise (HPE) [2]. Its modular design uses zero-carbon electricity and fan-less liquid cooling that reduce cooling power by up to **90 %** and cut carbon emissions by **72 %** [2].

- **Hardware.** The system uses **5 448 NVIDIA GH200 Grace-Hopper superchips** integrated into HPE's Cray EX architecture [3]. Each Grace-Hopper combines a 72-core Grace CPU with an H100-class GPU and high-bandwidth memory. The cluster delivers **≈21 AI exaflops** of performance [4].
- **Access for public sector.** The AIRR provides free access for UK researchers and start-ups via UK Research & Innovation (UKRI). Users apply through DSIT/UKRI allocation schemes. Inference workloads can run on the Grace-Hopper GPUs, which are optimised for mixed-precision matrix operations and memory-efficient LLM inference.
- **Role in AI inference.** The Grace-Hopper architecture integrates CPU and GPU on the same package, reducing data-transfer latency. This enables efficient inference of large language models (LLMs) and generative AI. The energy-efficient cooling and renewable power supply make it attractive for public-sector workloads that must meet sustainability criteria.

1.2 Cambridge Dawn (University of Cambridge – Cambridge Open Zettascale Lab)

- **Purpose & status.** Dawn is the UK's fastest AI supercomputer and the other half of the AIRR. The government invested **£36 million** to increase its capacity six-fold by spring 2026 [5]. Dawn has already supported over **350 projects** and provides free access to UK researchers and start-ups [5].
- **Hardware.** Dawn comprises **256 Dell PowerEdge XE9640 servers**. Each node contains **2 Intel Xeon Platinum 8468 processors (96 cores total), 1 TiB RAM and 4 Intel Data Center GPU Max 1550 (Ponte Vecchio) GPUs** [6]. The GPUs are connected by **Xe-Link** and quad-rail HDR200 InfiniBand. A near-term upgrade will add **AMD MI355X GPUs** integrated with StackHPC software [5].
- **Performance.** The system provides roughly **one exaflop** of AI performance and is the fastest AI supercomputer deployed in the UK [7].
- **Access.** Access is through the AIRR allocation route. Users manage projects via the Cambridge Research Computing Portal, which federates authentication and automatically provisions resources [8].
- **Inference relevance.** Ponte Vecchio GPUs contain 1024 compute units and have 128 GiB of HBM2e memory each [6]. These characteristics suit LLM inference (high memory bandwidth and large per-GPU memory). The planned addition of MI355X GPUs will diversify the technology landscape and broaden software support [5].

1.3 Sunrise – AI supercomputer for fusion power (University of Cambridge & UKAEA)

- **Purpose.** Announced in 2026, **Sunrise** is a £45 million AI supercomputer dedicated to accelerating nuclear-fusion research. Funded by the Department for Energy Security and Net Zero (DESNZ) and co-designed by Cambridge Research Computing, UK Atomic Energy Authority (UKAEA), Dell, AMD and StackHPC, Sunrise will support fusion plasma modelling and materials science [9].
- **Hardware & scale.** Sunrise is a **1.4 MW GPU-accelerated system**, making it one of the UK's largest AI supercomputers and the world's largest dedicated to fusion power [10]. Detailed hardware specifications have not yet been publicised; the project emphasises GPU acceleration and integration into DSIT's national AI service [11].
- **Access.** Sunrise will be integrated into the national AI ecosystem, enabling public-sector researchers to study plasma turbulence, materials development and fuel breeding for fusion energy [12].

1.4 Edinburgh International Data Facility (EIDF)

- **Overview.** The EIDF is a **state-of-the-art data analytics and storage service** built by EPCC to underpin the Data-Driven Innovation (DDI) programme [13]. It combines high-powered compute, flexible cloud services and specialised AI accelerators.
- **GPU service.** The EIDF GPU service offers a **containerised platform** accelerated by **16 NVIDIA H200, 136 NVIDIA H100 (80 GB) and 224 NVIDIA A100 GPUs** [14]. Users access resources through Kubernetes; a standard project namespace provides 100 CPU cores and 1 TiB memory [15].
- **Cerebras wafer-scale cluster.** EIDF hosts the largest Cerebras CS-3 cluster in Europe. **Four CS-3 Wafer-Scale Engines connected via a SwarmX switch** and a **MemoryX** streaming unit form a wafer-scale cluster optimised for training and inference of very large neural networks [16]. The system can train models up to **1 trillion parameters** and offers inference speeds of **>2 000 tokens s^{-1}**, up to **70× faster** than leading GPU solutions [17].
- **Usage & access.** The GPU and Cerebras services are available to University of Edinburgh researchers and DDI partners; external public-sector projects can apply via the EIDF service desk. The Cerebras platform simplifies LLM training and inference by avoiding distributed programming [18].

1.5 Planned National Supercomputer (EPCC – Edinburgh)

- **Overview.** The UK government is investing up to **£750 million** in a new national supercomputer at EPCC (Edinburgh). The system will come online in **2027**, replacing ARCHER2, and is expected to rank among the world's top five systems (of today's list) [1]. It will support AI research across climate, medicine and national security [19] [20].
- **Design.** Hardware specifications have not been released, but the compute roadmap notes that the machine will offer exascale performance, with a mix of CPU, GPU and other accelerators to support both training and inference. It will anchor a network of national supercomputing centres, improve energy efficiency and champion diverse suppliers [1].

2 Tier-0 & Tier-1 HPC centres accessible to public sector

2.1 ARCHER2 (EPCC – Edinburgh)

- **Role.** ARCHER2 is the UK's national supercomputing service (Tier-1) used for scientific research across academia and government. It is due to be succeeded by the new national supercomputer in 2027 but remains accessible.
- **Hardware.** ARCHER2 is an **HPE Cray EX system** providing **28 PFLOPS** peak performance. It contains **5 860 compute nodes** each with dual **AMD EPYC 7742 (64-core) processors**, giving a total of **750 080 CPU cores** [21]. Nodes have either **256 GiB (5 276 nodes) or 512 GiB (584 nodes)** of memory [22]. A small GPU development platform offers **4 nodes** with four AMD MI210 GPUs per node [22]. The Slingshot interconnect provides two 100 Gbps links per node [22].
- **Access.** UK researchers and public-sector partners apply via EPSRC's ARCHER2 allocation programme. While primarily CPU-focused, ARCHER2 can handle inference workloads that require large-scale parallelism (e.g., ensemble weather simulations or HPC-based ML pre-processing).

2.2 Mary Coombs (Hartree Centre – Daresbury Laboratory)

- **Purpose.** Launched in October 2025, **Mary Coombs** replaces the Scafell Pike system at the STFC Hartree Centre. It delivers **≈10 × the performance** of its predecessor while using **75 % less energy**

[23]. The facility is intended to accelerate industry-focused innovation and public-sector research in AI and data science [24].

- **Hardware.** Mary Coombs consists of **157 compute nodes powered by NVIDIA H100 GPUs**, forming a **24.41 PFLOP** GPU-based supercomputer [23] [24]. Nodes are integrated into a platform-as-a-service environment with flexible self-service and consultancy support [25].
- **Access.** The Hartree Centre offers the system as a PaaS to industry, academia and public-sector organisations, with on-demand GPU allocations for AI training and inference. The centre provides expertise in AI optimisation and visualisation.

2.3 NI-HPC / Kelvin-2 (Queen's University Belfast & Ulster University)

- **Overview.** Kelvin-2 is a Tier-2 HPC facility jointly run by Queen's University Belfast and Ulster University. The system was installed in 2022 and upgraded to include cutting-edge GPUs.
- **Hardware.** The cluster has **102 Dell PowerEdge R6525 nodes**, each with two AMD EPYC 7702 processors (128 cores) and 786 GB RAM [26]. There are **8 high-memory nodes** with **2 TB RAM**. GPU resources include **32 NVIDIA Tesla V100 GPUs**, **16 NVIDIA A100 GPUs**, **4 AMD MI300X GPUs**, **4 NVIDIA H100 GPUs** and **4 Intel Max 1100 GPUs** [26]. A **2 PB Lustre** file system supports high-bandwidth I/O and nodes are interconnected with EDR InfiniBand [26].
- **Access. 65 %** of resources are reserved for QUB and Ulster University; **35 %** is available free to researchers across the UK via EPSRC's HPC access process [27]. Kelvin-2 supports AI inference through its A100/H100/MI300X GPUs and is suitable for data-intensive applications such as neurotechnology and precision medicine.

2.4 Hartree GridPP (Worldwide LHC Computing Grid – UK Tier-1 & Tier-2 sites)

- **Overview. GridPP** provides distributed computing for the Large Hadron Collider and is part of the Worldwide LHC Computing Grid. It aggregates HPC clusters at UK universities and national labs. Up to **10 %** of resources (thousands of CPU cores and petabytes of storage) can be offered to new user communities, including industry [28].
- **Tier-1 centre.** The UK Tier-1 centre at Rutherford Appleton Laboratory (RAL) provides advanced processors, extensive storage and high-speed networking. It supports data processing for LHC experiments and offers secure, well-managed infrastructure [29].
- **Access.** Access is organised through GridPP and STFC. For AI inference, GridPP can supply large volumes of CPU cores and storage for embarrassingly parallel workloads or data processing stages in ML workflows.

2.5 IRIS federated infrastructure (STFC)

- **Overview. IRIS** is a community-led federation that coordinates digital research infrastructure across STFC laboratories and universities. It combines cloud services, HPC clusters, grid computing, tape archives and object storage [30].
- **Resources.** The federation integrates the STFC Cloud, the Arcus cloud at Cambridge, Imperial College Cloud and the Somerville Cloud at Edinburgh. It offers heterogeneous CPU and GPU resources, local clusters, and large-scale tape and S3 storage [31]. IRIS-CAST provides carbon-footprint tracking [32].
- **Access & allocation.** Resources are allocated through a **Resource Scrutiny and Allocation Process**, and identity management is federated across partners [31]. Public-sector projects can request

compute, including GPUs, via IRIS for AI inference workloads that need federated access or secure data handling.

2.6 JASMIN data analysis facility (NERC – CEDA)

- **Role.** JASMIN provides a unique combination of **tens of petabytes of storage** and compute resources for environmental science [33]. It supports interactive analysis, batch compute and a community cloud.
- **Upgrades.** In 2026, JASMIN added high-memory nodes with **6 TB RAM and 192 CPU cores each**, and introduced a GPU partition using **Multi-Instance GPU (MIG)** technology for ML workloads [34]. The MIG partition allows GPUs to be divided into smaller instances, improving efficiency for AI inference.
- **Access.** NERC scientists access JASMIN free of charge; other users can apply on a cost-recovery basis [35].

2.7 Met Office cloud supercomputer (Microsoft Azure)

- **Overview.** In 2025 the Met Office transitioned to a **cloud-based supercomputer** delivered as a managed service by Microsoft. It is the first such system dedicated to weather and climate science [36]. The system combines the Met Office's modelling expertise with Azure's infrastructure and forms part of the UK's national capability.
- **Capacity.** Microsoft states that the system comprises **≈1.8 million CPU cores** and is capable of **≈60 quadrillion calculations per second** [37]. This places it among the most powerful CPU-based supercomputers globally and the most advanced dedicated to weather forecasting [37].
- **Benefits & access.** The cloud model provides flexible capacity scaling and resilience. Public-sector users (e.g., environment agencies) benefit from improved weather forecasts and can collaborate with the Met Office on AI-enhanced modelling. While direct access for inference is limited, the system demonstrates how cloud supercomputing can complement on-premise AI inference.

2.8 AWE high-performance computing

- **Context.** The Atomic Weapons Establishment (AWE) operates one of the UK's most powerful supercomputing programmes to support national security and safety analysis. Details are not publicly disclosed, but the UK compute roadmap lists AWE's HPC centre as part of the public research and innovation compute ecosystem [38].
- **Relevance.** While access is restricted and primarily for defence, the presence of such capacity underscores the UK's overall compute capability. Collaboration with AWE in AI algorithms (e.g., nuclear safety modelling) may eventually translate into broader public-sector applications.

3 DiRAC national services – extreme scaling, data-intensive & memory-intensive

The **Distributed Research using Advanced Computing (DiRAC)** programme provides specialised HPC services for particle physics, astrophysics and cosmology. Although designed for fundamental science, the systems are accessible via STFC and can be used for AI inference requiring huge datasets or parallel GPU performance.

DiRAC service	Location & owner	Hardware & capacity	Relevance for AI inference
Extreme Scaling (Tursa)	EPCC (Edinburgh)	64 GPU nodes each with **4 NVIDIA A100 80 GB GPUs**, 114 GPU nodes with **4 NVIDIA A100 40 GB GPUs**, 6 CPU nodes with AMD EPYC 7H12, giving **3 504 CPU cores**, **291 840 Tensor cores**, **>180 TB RAM** and **5.1 PB Lustre** storage [39] [40].	Water-cooled racks and non-blocking HDR InfiniBand enable large-scale GPU inference or training; particularly suited for particle-physics simulations and large-language-model inference.
Data-Intensive	Cambridge Service for Data-Driven Discovery (CSD3)	30 412 CPU cores and 746 496 GPU cores, 157 TB memory, 8 PB storage, HDR interconnect with 3:1 blocking [41].	Supports data-rich workloads; can process large datasets for ML.
Data-Intensive (Leicester)	University of Leicester	40 288 CPU cores (1 216 chips), 177 TB memory, 7 PB storage [41].	Provides high I/O throughput; suitable for AI workflows needing big-data pre-processing.
Memory-Intensive	Durham University	80 240 CPU cores, 731 TB memory, 17.1 PB storage with non-blocking HDR InfiniBand [41].	Offers large per-node memory for simulations and deep neural network training with huge parameter counts.
Access	Projects apply via the DiRAC Resource Allocation Committee; commercial and trial access are possible [40].	AI inference tasks requiring extreme parallelism or memory can be accommodated when they align with DiRAC's mission.	

4 Regional and Tier-2 HPC systems

4.1 Cirrus (EPCC – Edinburgh)

- **Overview.** Cirrus is a general-purpose Tier-2 system for computational science. In 2025 it was upgraded to an **HPE Cray EX4000** cluster with **256 CPU nodes** using dual **AMD EPYC 9825** processors (144 cores per node) and a total of **73 728 CPU cores** [42]. Nodes are water-cooled and connected via **Slingshot 11 (200 Gbps)**; there are **192 nodes with 768 GB memory** and **64 high-memory nodes with 1.5 TB** [42] [43]. Storage comprises **1 PB Ceph** and **1 PB Lustre** [42].

- **Access & AI relevance.** Cirrus has no GPUs, making it suitable for CPU-bound inference (e.g., inference on CPU-optimised models, data pre-processing or ensemble evaluation). Access is via EPSRC and industry collaborations; it can also be used through the EIDF portal.

4.2 Bede (N8 Research Partnership – Durham University)

- **Overview.** Bede serves researchers in the N8 universities. It uses **32 IBM Power9 dual-CPU nodes**, each with **4 NVIDIA V100 GPUs**, and includes **4 nodes with NVIDIA T4** GPUs for inference [44]. Since March 2024 Bede also offers **3 NVIDIA Grace-Hopper nodes** for pilot projects [44].
- **Access.** N8 researchers can access Bede without charge; external researchers may apply via EPSRC. Bede's GPU nodes and T4 inference nodes support deep-learning workloads and AI inference at small to medium scale.

4.3 Materials & Molecular Modelling Hub (Young facility – London)

- **Overview.** The **Young** supercomputer supports materials and molecular modelling. It consists of **23 040 CPU cores** on HPE dual 20-core Intel nodes and **6 nodes with 8 NVIDIA A100 GPUs** and dual AMD EPYC 7543 CPUs [45]. **70 %** of capacity is reserved for MMM Hub partners; **30 %** is free to UK researchers [45].
- **Relevance.** Although targeted at chemistry, the A100 nodes can run AI inference for materials science. Projects apply via the MMM Hub.

4.4 Cambridge Service for Data-Driven Discovery (CSD3) CPU & GPU clusters

CSD3 hosts several sub-clusters; the table summarises hardware relevant to AI inference.

Cluster	Nodes & processors	GPU/accelerators	Interconnect & memory	Notes
Ampere GPU (Wilkes3) [46]	**90 Dell PowerEdge XE8545** servers; each node has **2× AMD EPYC 7763 (128 cores)**, **1 TB RAM** [46].	**4× NVIDIA A100-SXM 80 GB GPUs** per node [46].	Dual-rail **Mellanox HDR200** InfiniBand; each A100 has 6 912 FP32 CUDA cores [47].	Known as **Wilkes3**; available via CSD3 GPU partition.
Cascade Lake CPU [48]	Nodes named cpu-p-[1-672]; each has **56 CPU cores** and **192 GiB RAM** (or **384 GiB** for high-memory nodes) [49].	None.	**Mellanox HDR100** InfiniBand [50].	For CPU-bound workloads.
Ice Lake CPU [51]	Nodes cpu-q-[1-544]; each has **76 CPU cores** and **256 GiB RAM** (512 GiB for high-memory nodes) [52].	None.	**Mellanox HDR200** InfiniBand [53].	Supports data processing and CPU inference.

Cluster	Nodes & processors	GPU/accelerators	Interconnect & memory	Notes
Sapphire Rapids CPU [54]	Nodes cpu-r-[1-112]; each has **112 CPU cores** and **512 GB RAM** [55].	None.	**Mellanox NDR200** InfiniBand [56].	Newest CPU nodes for general HPC.
Sapphire Rapids with High-Bandwidth Memory (HBM) [57]	Nodes cpu-r-hbm-[1-84]; each provides **112 CPU cores** and **1 TiB RAM** (1.15 TiB on flat partition) [58].	None.	Mellanox NDR200; runs **Rocky Linux 9** [59].	Exclusive to UKAEA projects.

- **Access & AI relevance.** GPU nodes (A100) support deep-learning training and inference; CPU nodes provide pre- and post-processing. CSD3 is accessible to academic and industrial users via EPSRC allocations. Dawn and the EIDF GPU service are integrated into CSD3, providing a continuum from local to national AI compute.

4.5 JADE II (Joint Academic Data Science Endeavour – decommissioning)

- **Overview.** JADE II was a national Tier-2 GPU cluster supporting machine-learning research. It has been deprecated and will be decommissioned in early 2025 [60].
- **Hardware.** The cluster comprised **63 DGX MAX-Q servers**, each with **8 NVIDIA V100 (32 GB) GPUs** and **512 GB RAM**, connected by EDR InfiniBand [61]. Shared storage included **70 TB NVMe** and **1 PB Lustre** [61].
- **Access & relevance.** Members of the University of Sheffield could access JADE II for free; other UK researchers used EPSRC allocations. Despite its decommissioning, JADE II set a precedent for national GPU facilities.

5 Other computing resources listed in the UK compute roadmap

The compute roadmap annex identifies additional facilities that may support AI inference [38]. Some are primarily research clusters but can be leveraged for public-sector AI.

Facility	Summary & relevance
GridPP Tier-2 sites	Regional grid clusters across the UK provide thousands of CPU cores and petabytes of storage for parallelisable tasks [28]. These are suitable for inference pipelines requiring large-scale data processing.
Imperial College, Cambridge Arcus, Somerville Cloud (part of IRIS)	University-operated cloud platforms with heterogeneous CPUs and GPUs [31]. Support containerised AI workloads and provide federated identity management.
University of Durham (Bede)	See Section 4.2. Supports GPU-accelerated inference and includes Grace-Hopper nodes for exploring CPU-GPU integration [44].

Facility	Summary & relevance
Bristol Isambard 3	CPU-based supercomputer using **NVIDIA Grace CPU Superchips** with **>55 000 cores**, offering six-times the performance of Isambard 2 [62]. It supports energy-efficient HPC and is co-located with Isambard-AI.
EIDF Data Science Cloud	Private cloud built on OpenStack and Ceph; provides flexible VM provisioning, GPU hypervisors and access to the **Cerebras CS-3 wafer-scale cluster** [63] [64]. Supports novel AI development and secure research computing.
Met Office HPC facility (operational 2025)	See Section 2.7. Offers 1.8 million CPU cores and 60 quadrillion operations s^{-1} for weather and climate modelling [37].
National infrastructure at AWE	High-performance systems used for national security; details undisclosed but included in the compute ecosystem [38].

6 Leveraging UK compute capacity for AI inference

6.1 Federated access and cloud integration

The UK's compute landscape is fragmented across universities, national labs and cloud providers. Initiatives like **AIRR**, **IRIS**, the **EIDF portal** and the **Cambridge Research Computing access portal** provide federated identity management and automated resource provisioning [8] [31]. Public-sector organisations should register through these portals to request AI inference capacity.

6.2 Balancing training and inference

The compute roadmap emphasises a balance between training and inference to maximise utility [1]. For inference workloads, high-bandwidth memory and large per-GPU memory (H100, Grace-Hopper, Ponte Vecchio) are critical. Facilities such as Isambard-AI, Dawn and the EIDF GPU service offer this capacity. When inference is CPU-bounded or memory-intensive (e.g., vector databases, Monte-Carlo inference), CPU-based clusters like ARCHER2, Cirrus or the Met Office system provide scale.

6.3 Energy efficiency and sustainability

Public-sector projects must consider energy use. Many UK systems incorporate energy-saving features:

- **Isambard-AI** uses renewable electricity and liquid cooling to cut cooling power by **up to 90 %** [2].
- **Mary Coombs** delivers **10× performance** while reducing energy consumption by **75 %** compared with its predecessor [23].
- The **Met Office cloud supercomputer** is powered by 100 % sustainable energy and aims to save thousands of tonnes of CO_2 in its first year [65].

6.4 Public-sector use cases

Public-sector organisations have already leveraged these facilities:

- **Climate and environmental forecasting.** The **Dawn** supercomputer runs *IceNet*, a sea-ice forecasting model that uses AI to deliver seasonal forecasts; once trained on Dawn, the model can run on a laptop, enabling rapid inference in the field [66].
- **Medical diagnostics.** Cambridge researchers use Dawn to train a deep-learning tool for kidney-cancer screening; after training, the model offers fast and accurate inference for CT scans [67].
- **Environmental data analysis.** JASMIN's high-memory and GPU-MIG nodes allow scientists to analyse petabytes of earth-observation data and run machine-learning models for environmental risk assessment [34].

Conclusion

The UK hosts a diverse and expanding set of computing resources that position it to be a **leader in AI inference** even if model development occurs elsewhere. National AI supercomputers like **Isambard-AI**, **Dawn** and **Sunrise**, combined with federated infrastructures such as **IRIS** and **EIDF**, give public-sector researchers access to GPUs, wafer-scale accelerators, CPU supercomputers and flexible clouds. Tier-1 and Tier-2 systems (ARCHER2, Mary Coombs, Kelvin-2, Cirrus, Bede and CSD3) provide additional capacity for data processing, CPU-based inference and hybrid workflows. The compute roadmap's commitment to expanding this ecosystem to **420 AI-exaflops** by 2030 [1] ensures that the UK's public sector can meet the growing demand for AI inference, support critical applications from climate forecasting to healthcare, and maintain resilience and sustainability in the face of rapid technological change.

---

[1] [38] UK Compute Roadmap - GOV.UK
https://www.gov.uk/government/publications/uk-compute-roadmap/uk-compute-roadmap

[2] [4] 2025: Isambard-AI launches July 2025 | Research | University of Bristol
https://www.bristol.ac.uk/research/centres/bristol-supercomputing/articles/2025/isambard-ai-launches-july-2025.html

[3] Isambard-AI and Dawn AIRR supercomputers: Innovator route – UKRI
https://www.ukri.org/opportunity/isambard-ai-and-dawn-airr-supercomputers-innovator-route/

[5] Cambridge supercomputer set to get 6 times more powerful as government backs British AI innovation - GOV.UK
https://www.gov.uk/government/news/cambridge-supercomputer-set-to-get-6-times-more-powerful-as-government-backs-british-ai-innovation

[6] [8] Dawn - Intel GPU (PVC) Nodes — CSD3 1.0 documentation
https://docs.hpc.cam.ac.uk/user-guide/pvc.html

[7] Dawn
https://www.hpc.cam.ac.uk/d-w-n

[9] [10] [11] [12] £45 million for the UK's first AI supercomputer dedicated to fusion power | University of Cambridge
https://www.cam.ac.uk/news/ps45-million-for-the-uks-first-ai-supercomputer-dedicated-to-fusion-power

13 63 64 EPCC and StackHPC deliver a highly-available data science cloud for EIDF | EPCC
https://www.epcc.ed.ac.uk/whats-happening/articles/epcc-and-stackhpc-deliver-highly-available-data-science-cloud-eidf

14 15 EIDF GPU Service | Edinburgh International Data Facility | Edinburgh International Data Facility
https://edinburgh-international-data-facility.ed.ac.uk/services/computing/gpu-service

16 EIDF Cerebras | Edinburgh International Data Facility | Edinburgh International Data Facility
https://edinburgh-international-data-facility.ed.ac.uk/services/computing/cerebras-cs

17 18 Cerebras Expands AI Training & Inference Cluster in Europe at the University of Edinburgh to Accelerate AI Research and UK's AI Capabilities
https://www.cerebras.ai/press-release/cerebras-expands-ai-training-and-inference-cluster-in-europe-at-the-university-of-edinburgh-to

19 University set to host £750m national supercomputer - DDI
https://ddi.ac.uk/university-set-to-host-750m-national-supercomputer/

20 University named UK's first National Supercomputing Centre | News | The University of Edinburgh
https://www.ed.ac.uk/news/university-named-uks-first-national-supercomputing-centre

21 22 ARCHER2 Hardware & Software
https://www.archer2.ac.uk/about/hardware.html

23 STFC launches powerful new AI supercomputer for industry - Hartree Centre
https://www.hartree.stfc.ac.uk/news/2025/10/20/stfc-launches-powerful-new-ai-supercomputer-for-industry/

24 25 Mary Coombs - Hartree Centre
https://www.hartree.stfc.ac.uk/technologies/mary-coombs/

26 Kelvin2 | Northern Ireland High Performance Computing
https://ni-hpc.ac.uk/Kelvin2/

27 About | Northern Ireland High Performance Computing
https://ni-hpc.ac.uk/about/

28 For Industry – GridPP
https://www.gridpp.ac.uk/users/industry/

29 The UK Tier-1 Centre – GridPP
https://www.gridpp.ac.uk/infrastructure/the-uk-tier-1-centre/

30 31 32 IRISbookletFinal.pdf
https://www.sc.stfc.ac.uk/wp-content/uploads/2025/03/IRISbookletFinal.pdf

33 35 JASMIN | NERC Environmental Data Service
https://eds.ukri.org/services/jasmin

34 JASMIN updates March 2026 - Centre for Environmental Data Analysis
https://www.ceda.ac.uk/news/updates/2026/2026-03-11-jasmin-update-march-2026/

36 65 Our supercomputer for weather and climate forecasting - Met Office
https://www.metoffice.gov.uk/about-us/who-we-are/innovation/supercomputer

37 A year of cloud-powered weather forecasting: how the UK's national capability is evolving - Met Office
https://www.metoffice.gov.uk/blog/2026/a-year-in-the-cloud-the-met-office-supercomputer

[39] [40] DiRAC: Tursa (GPU) | EPCC
https://www.epcc.ed.ac.uk/hpc-services/dirac-tursa-gpu

[41] Our Services – DiRAC High Performance Computing Facility
https://dirac.ac.uk/our-services/

[42] Cirrus | EPCC
https://www.epcc.ed.ac.uk/hpc-services/cirrus

[43] Cirrus Hardware and Software
https://www.cirrus.ac.uk/about/hardware-software/

[44] Bede Supercomputer
https://n8cir.org.uk/bede/

[45] mmmhub.ac.uk
https://mmmhub.ac.uk/

[46] [47] Ampere GPU Nodes — CSD3 1.0 documentation
https://docs.hpc.cam.ac.uk/user-guide/a100.html

[48] [49] [50] Cascade Lake Nodes — CSD3 1.0 documentation
https://docs.hpc.cam.ac.uk/user-guide/cclake.html

[51] [52] [53] Ice Lake Nodes — CSD3 1.0 documentation
https://docs.hpc.cam.ac.uk/user-guide/icelake.html

[54] [55] [56] Sapphire Rapid Nodes — CSD3 1.0 documentation
https://docs.hpc.cam.ac.uk/user-guide/sapphire.html

[57] [58] [59] Sapphire Rapid Nodes with High Bandwidth Memory — CSD3 1.0 documentation
https://docs.hpc.cam.ac.uk/user-guide/sapphire-hbm.html

[60] [61] JADE II (Tier 2 GPU cluster) — Sheffield HPC Documentation
https://docs.hpc.shef.ac.uk/en/latest/other-uk-hpc-resources/jade2.html

[62] 2024: Cutting-edge research boost as GW4 Isambard 3 supercomputer goes online | Cabot Institute for the Environment | University of Bristol
https://www.bristol.ac.uk/cabot/news/2024/isambard.html

[66] [67] Under the bonnet at Dawn, the UK's fastest AI supercomputer
https://www.cam.ac.uk/stories/under-the-bonnet-at-AI-supercomputer-Dawn